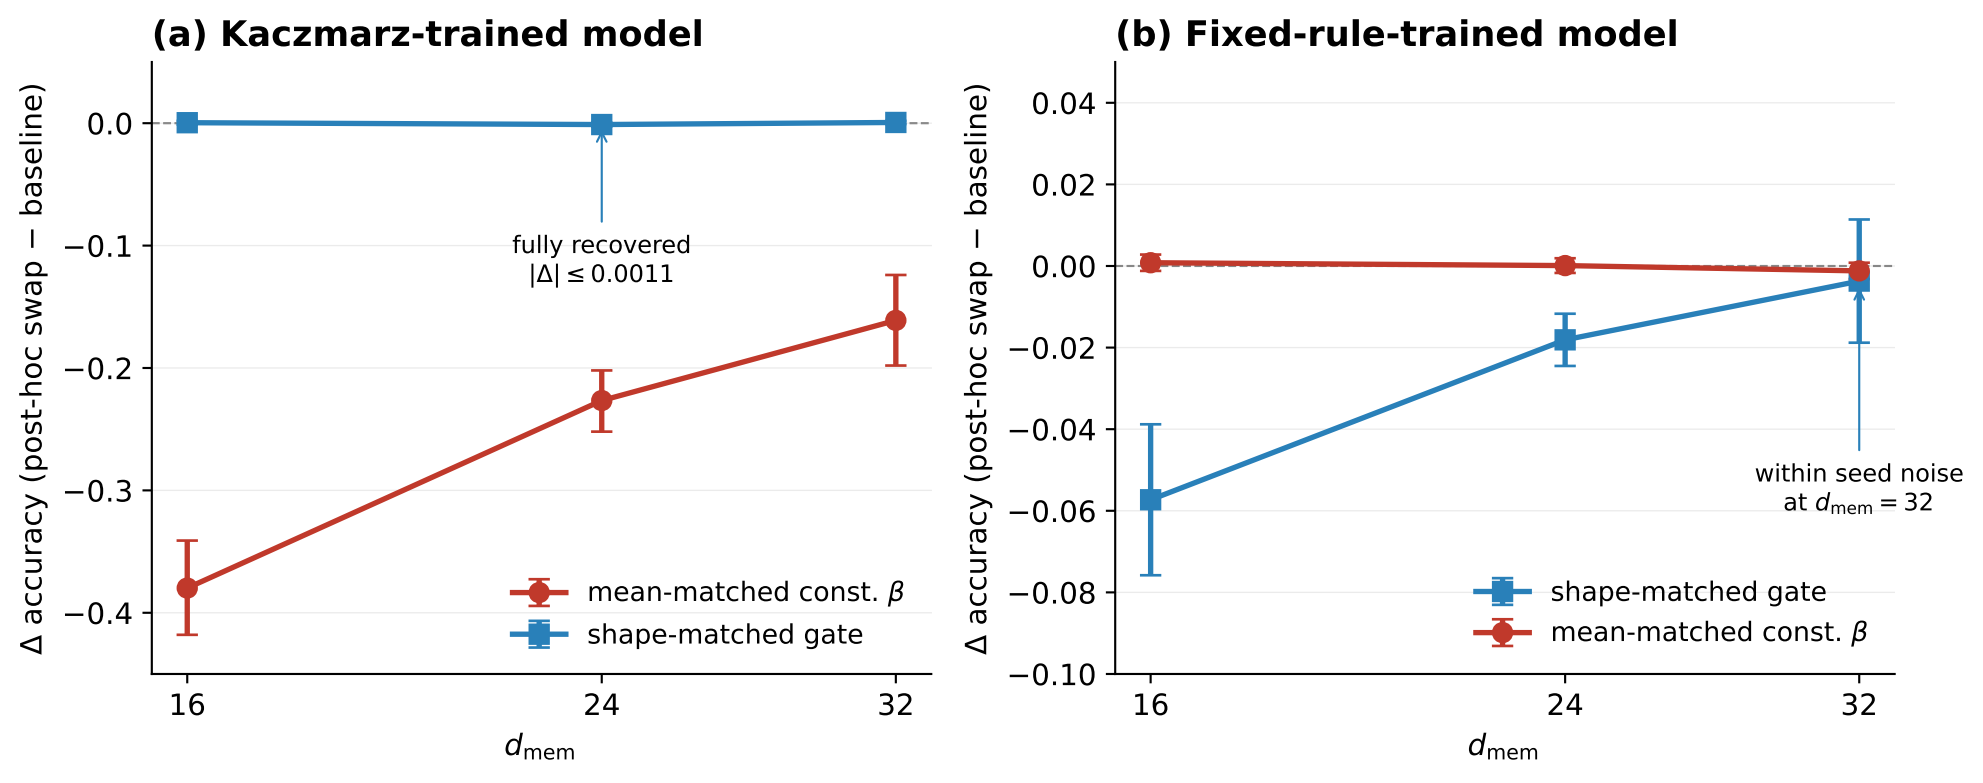
(a) Kaczmarz-trained model
Δ accuracy (post-hoc swap − baseline)
0.0
−0.1
−0.2
−0.3
−0.4
fully recovered
|Δ| ≤ 0.0011
mean-matched const. β
shape-matched gate
16
24
32
d_mem
(b) Fixed-rule-trained model
Δ accuracy (post-hoc swap − baseline)
0.04
0.02
0.00
−0.02
−0.04
−0.06
−0.08
−0.10
within seed noise
at d_mem = 32
shape-matched gate
mean-matched const. β
16
24
32
d_mem